

A2 — Researcher and Doctoral-Student Survey

Self-reported AI practices, exposure, knowledge, and context among doctoral students, so leadership's documentary account can be compared with respondents' accounts.

About this instrument

Purpose	To collect self-reported practices, exposure, knowledge, and context, making discrepancies with the institutional self-assessment (A1) visible and measurable.
Who completes it	All doctoral students at every stage of their program. Invite the full eligible population rather than drawing a sample, and report the achieved response rate and respondent profile.
Time required	10 minutes for the 42 closed items and one short-text item, while the two optional open-comment items add two to three minutes. The 24-item short form (★) takes about five minutes.
Related report sections	§1.1, §1.2, §1.3, §2.1, §2.2, §2.3, §2.4, §2.5, §3.1, §3.3, §4.2, §5.3, §5.5, Appendix A, Appendix C, Appendix G.
Grid cells touched	Nine of the twenty: D1-people, D1-process, D2-people, D2-process, D3-people, D3-systems, D3-process, D4-people, and D4-process.
Cells not assessed	The five policy-axis cells. Whether a policy exists is a documentary question answered by whether A1 can produce the document. A2 records whether respondents report that the policy reached them in practice.
Companion instruments	A1 (institutional self-assessment), A3 (supervisor survey, with twelve items addressing related topics), A4 (gap-analysis worksheet, where the three sources are reconciled).
What a respondent sees	The two statements below, followed by Sections 1 to 10. Everything after Section 10 is for the analyst and contains the answer key, so it should not be printed or shown to respondents.

Table 1. Key specifications, purpose, and administrative details of the survey instrument.

Dimension key. D1 human-in-the-loop discipline · D2 responsible use in practice · D3 tooling that promotes responsible use · D4 AI-literate humans · D5 institutional benchmarking grid. Each is scored across four axes (policy, people, systems, process), giving the twenty cells.

Why this instrument exists. Appendix A of [Responsible AI in Academic Research: A Competency Framework for Research Training](#) names "a graduate-research survey instrument" among the evidence an institution needs to reach *leading* on Dimension 4. This is that instrument.

The response scale

Scale A. Yes · Partial · No · I don't know (used for Q10, Q12, and Q39, with Q20, Q25, and Q33 using slight variants). *Partial* means true in part, sometimes, or for one piece of work but not others. *I don't know* is a valid and important answer that is reported separately and never added to *No*. An institution whose graduate students cannot find its rule has a very different problem from an institution that has no rule at all.

Scale B (how often). Q4 to Q7: Never · A few times in the past three months · About once a month · About once a week · Most days

Scale C (how much). Q16: All · Most · About half · A few · None · I have not done this

Check-all lists. Q9, Q13, Q21, Q23, Q36, and Q41, marked (*tick all that apply*).

Q35 and Q40 ask what you think rather than what you did, so that beliefs can be compared with self-reported practices. Every other item asks about self-reported practice, exposure, knowledge, or what you have been told.

Before you begin: how your answers are used

Getting honest answers depends on making this clear.

Anonymous. The survey does not ask for your name, student number, the name of your advisor or supervisor, department, or year of entry. Please do not write those in the open comments.

Reported in groups of at least ten. Smaller groups are combined before anything is reported, and results for a school or college will appear only when it has at least ten responses.

Open comments are edited before reporting. Any detail that could identify you, your advisor or supervisor, or a specific project is removed.

Your answers cannot be used against you. They form no part of any progress review or milestone, your examination, or any research-misconduct or academic-integrity process. [INSTITUTION NAME] has confirmed this in writing in [NAME OF CONFIRMING DOCUMENT, DATE], readable at [URL].

Some questions ask about things you may have been told not to do. Please answer them honestly, as covered in the assurance above. If that protection is not formally in place at your institution, you should not run this survey.

Where your answers are held. Responses are collected by [SURVEY PLATFORM], stored in [COUNTRY OR REGION], held by [OFFICE], and will be destroyed on [DATE]. Only [NAMED ROLE] and [NAMED ROLE] see raw responses.

Voluntary. You may skip any question and stop at any time. Partial responses are still counted.

What happens next. A summary goes to [COMMITTEE] on [DATE] and the same summary goes to every doctoral student. If the results are not shared with you, this survey has not been run properly.

Drafting note for the institution. Secure the no-detriment confirmation in writing from your research integrity office and graduate school before the survey opens, and name the document above. Q36 to Q39 ask respondents to describe how they handle unpublished and participant data. Without a clear written amnesty, those four items will be underreported, and underreporting there is worse than not asking at all because it gives a false sense of reassurance. If [SURVEY PLATFORM] moves responses outside your jurisdiction, make sure to state that clearly above.

A note on words

Doctoral student is used throughout and covers every stage of enrollment: graduate student, PhD student, doctoral candidate, higher degree by research (HDR) candidate, research student, postgraduate researcher, doctoral researcher, *Doktorand*, and *doctorant* all refer to the same person.

Thesis means the thesis or dissertation (the major written work you submit at the end), whichever term your institution uses. **Examination** covers however that work is assessed, including examiners' written reports, an oral examination, or a defense.

AI tool means any tool that generates text, code, images, summaries, transcripts, translations, or citations for you, whether or not it is marketed as artificial intelligence.

Supervisor and **advisor** mean the same person here: your principal supervisor, main advisor, dissertation advisor, *Doktorvater* or *Doktormutter*, or whoever signs off on your progress.

Items are kept short and plain because the framework behind this survey covers thirty-eight universities across fifteen countries and jurisdictions (the report §1.3), and many respondents do not speak English as their first language. If you translate this survey, please translate both the question stems and response options together, and keep the original option order.

Key to the markings

- ★: Part of the 24-item short form.
- ⇔ **PAR-XXXX**: A topic-link key. A3 uses a subset of these keys to connect related student and supervisor items. The key supports joining the two files, but it does not establish that the items are directly comparable. Keys marked **(core)** identify priority topics when an institution creates a shorter local A3.
- **Cells**: The maturity-grid cells that this item informs.
- **Evidence anchor**: The institutional record used to interpret or check the aggregate answer. This is for the analyst and is **not shown to respondents**. A survey answer on its own is not auditor-grade evidence, so the anchor identifies the relevant document or record.

Section 1: About you (3 items)

Q1 ★ Where are you in your doctorate?

- ☐ First year, before my first formal progress review or milestone ☐ Past that first review, in the middle years ☐ Final year ☐ Submitted, waiting for examination ☐ Other

Cells: D1-process · Evidence anchor: Enrollment records, used to check response rates by stage against the overall student population.

Q2 ★ Which best describes your field?

- ☐ Humanities ☐ Social sciences ☐ Business, law, economics ☐ Life and health sciences ☐ Physical sciences, engineering, computing ☐ Creative and performing arts ☐ Other

Cells: D4-process · Evidence anchor: Enrollment counts by school or college. Report by these seven broad groups only, because department-level reporting risks identifying individual respondents.

Q3 Is English your first language?

- ☐ Yes ☐ No ☐ One of several first languages

Cells: D3-systems, D4-people · Evidence anchor: None at the institutional level. This item exists because transcription tools and AI-detection tools carry documented bias against speakers whose first language is not English (the report Appendix C, Classes 10 and 11), so all other survey results should be analyzed across this split.

Section 2: Which AI tools you use (6 items)

Q4 to Q7 use **Scale B** and cover the four modes of AI use identified in the report: search, co-author, validator, and tutor.

Q4 ★ ⇌ PAR-MODE-SEARCH In the past three months, how often have you used an AI tool to **find or scope literature?**

- ☐ Never ☐ A few times in the past three months ☐ About once a month ☐
About once a week ☐ Most days

Cells: D2-people, D3-systems · Evidence anchor: Institutional license and seat-usage reports for AI search and retrieval tools, as well as library discovery service logs.

Q5 ★ ⇌ PAR-MODE-COAUTHOR In the past three months, how often have you used an AI tool to **draft, rewrite, shorten or polish your own writing?**

- ☐ Never ☐ A few times in the past three months ☐ About once a month ☐
About once a week ☐ Most days

Cells: D2-people, D2-process · Evidence anchor: Thesis submission form disclosure fields, and disclosure statements on submitted manuscripts held by the research office.

Q6 ★ ⇌ PAR-MODE-VALIDATOR In the past three months, how often have you used an AI tool to **check, critique or find errors in your own argument, method or results?**

- ☐ Never ☐ A few times in the past three months ☐ About once a month ☐
About once a week ☐ Most days

Cells: D2-people, D4-people · Evidence anchor: None directly available. This is the mode of AI use that is least visible to the institution and least covered by existing disclosure rules, which is why this survey asks students directly.

Q7 ★ ⇌ PAR-MODE-TUTOR In the past three months, how often have you used an AI tool to **explain a concept, a method or a piece of statistics to you?**

- ☐ Never ☐ A few times in the past three months ☐ About once a month ☐
About once a week ☐ Most days

Cells: D4-people, D4-process · Evidence anchor: Enrollment and completion records for institutional methods training. High usage of tutor mode paired with low methods training enrollment is a curriculum signal rather than a compliance issue.

Q8 ★ Which account do you mostly use for research work?

- ☐ An account [INSTITUTION NAME] provides ☐ A free personal account ☐ A
paid personal account I pay for myself ☐ An account provided by a
collaborator or an external partner ☐ A tool that runs on my own computer ☐
I do not use AI tools for research

Cells: D3-systems, D3-process · Evidence anchor: The institutional AI tenant's seat and license reports, IT expenditure records, and service-desk tickets concerning unsupported tools.

Q9 Which of these have you used for research work in the past three months? (tick all that apply)

- ☐ Writing or fixing code, or building an analysis pipeline ☐ Screening titles and abstracts for a review ☐ Transcribing interviews, focus groups or fieldnotes ☐ Translating text ☐ Making images or diagrams for talks or teaching ☐ Making images or figures that present data ☐ AI features inside tools I do not think of as AI tools: search-engine answers, reference-manager summaries, meeting transcription, writing assistants in email or word processors ☐ None of these

Cells: D3-systems, D2-people · Evidence anchor: The institutional tool register and procurement records. Tools selected here that do not appear on the official register represent unapproved shadow IT use. Look at the data-figure option first, because academic publishers are broadly converging against any AI-generated data figures (the report Appendix C, Class 9).

Section 3: What you have been told (6 items)

Q10 ★ ⇌ PAR-POLICY-KNOWN (core) Does [INSTITUTION NAME] have a written rule or policy about using AI in doctoral research?

- ☐ Yes ☐ Partial ☐ No ☐ I don't know

Cells: D2-people, D1-people · Evidence anchor: The policy document itself, along with its original publication date and last review date. Any gap between the proportion of students answering Yes and the actual existence of the document indicates a communication breakdown rather than an absence of policy.

Q11 ⇌ PAR-POLICY-LOCATE (core) If yes or partial: what is it called, or where is it? (Short text: a title, a web address, or a description such as "a page on the graduate school site" is enough. Write "cannot find it" if that is the honest answer.)

Response:

Cells: D2-people · Evidence anchor: The actual policy title and official URL. Analysts should code each response as correct, wrong document, or vague. A respondent who can name the correct document has demonstrably located it.

Q12 ★ ⇌ PAR-TASK-RULES (core) Have you been told **in writing** which research tasks you may use AI for and which you may not?

- ☐ Yes ☐ Partial ☐ No ☐ I don't know

Cells: D1-people, D1-process · Evidence anchor: The institution's published labor-versus-judgment guidelines (the report §2.1, Appendix G), the orientation packet, and the supervision agreement template.

Q13 ★ Where did you learn what the rules are? (tick all that apply)

- ☐ Induction or orientation ☐ A methods or research-training course ☐ My advisor or supervisor ☐ The thesis submission form ☐ An email or a web page from my school or college, or from the graduate school ☐ Other doctoral students ☐ The AI tool's own terms or warnings ☐ A journal or funder, not my institution ☐ I have not learned them

Cells: D2-people, D4-process · Evidence anchor: Orientation agendas and attendance records, methods course outlines, and graduate school communication logs. When "other doctoral students" ranks higher

than formal institutional channels, it shows that policies are being communicated primarily through informal word of mouth.

Q14 ★ ⇌ PAR-THESIS-DECLARE (core) Do you know whether you must declare AI use when you submit your thesis?

- ☐ Yes, and I know what I must declare ☐ Yes, but I do not know what I must declare ☐ No, there is no requirement ☐ I don't know

Cells: D1-process, D2-process · Evidence anchor: The thesis declaration form or affidavit and its AI disclosure clause, along with the completion rate for that clause over the last two submission cycles.

Q15 Do you know what the rules are for using AI in **grant or scholarship applications**, and in **manuscripts you submit to journals**?

- ☐ Yes, both ☐ One but not the other ☐ No ☐ I don't know ☐ Not relevant to me yet

Cells: D2-people, D2-process · Evidence anchor: Funder policy summaries held by the research office, as well as publisher guidance issued by the library or research office. External funder and publisher rules bind researchers directly, regardless of internal institutional guidance.

Section 4: Checking what the tool gives you (6 items)

Q16 ★ ⇌ PAR-CITE-CHECK (core) In your most recent literature search that used an AI tool: how many of the references it gave you did you open and read at the original source?

- ☐ All ☐ Most ☐ About half ☐ A few ☐ None ☐ I have not done this

Cells: D4-people, D2-process · Evidence anchor: Reference manager libraries attached to progress reviews or mid-program submissions, as well as attendance records for library research instruction.

Q17 ⇌ PAR-FAKE-REF-SEEN Have you ever found that a reference an AI tool gave you did not exist, or did not say what the tool claimed it said?

- ☐ Yes ☐ No ☐ I have not checked closely enough to know

Cells: D4-people · Evidence anchor: None at the institutional level, because this measures self-reported incidence. The third option ("I have not checked closely enough to know") provides an important finding. Measured citation fabrication rates range from 18 percent to 55 percent depending on the AI model (the report §2.4, competency 1). A low "yes" rate is not evidence that the tools are infallible. Interpret it alongside Q16 and Q18 before drawing conclusions about reference verification.

Q18 ★ ⇌ PAR-CITE-PROBE (core) Which **one** of these confirms that a reference is real and says what an AI tool claims it says?

- ☐ The tool gave a DOI or a link ☐ The reference appears in the tool's list of sources ☐ You opened the source and read the passage that supports the claim ☐ A second AI tool agreed that the reference is correct

Cells: D4-people, D4-process · Evidence anchor: Methods curriculum learning outcomes related to citation verification. The correct answer is provided in the administration notes below.

Q19 ⇌ PAR-MODEL-HETERO (core) When you use an AI tool to challenge or check your own argument, do you use a tool based on a **different underlying model family** from the one that helped you produce the argument?

☐ Yes, always ☐ Sometimes ☐ No ☐ I had not thought about this ☐ I do not use AI to check my arguments

A different product or company may still use the same underlying model. Check the model name and provider rather than relying on the product brand.

Cells: D4-people · Evidence anchor: Methods curriculum coverage of model heterogeneity in adversarial review (the report §2.4, competency 4). Report "I had not thought about this" as its own distinct category.

Q20 ★ Have you ever included AI output in a thesis chapter, a manuscript or an application when you could not check it yourself?

☐ Yes ☐ Partial (I checked some of it) ☐ No ☐ I would rather not answer

Cells: D1-people, D2-process · Evidence anchor: None, by design. This item measures the core threshold of the competency framework (specifically whether an AI tool is generating research content that the researcher cannot independently verify, the report §2.2, co-author mode), which is why an explicit written amnesty is essential.

Q21 The last time an AI tool gave you an answer that sounded confident and turned out to be wrong, what did you do? (tick all that apply)

☐ Checked it against the source and corrected it ☐ Asked the same tool again
☐ Asked a different tool ☐ Asked my advisor, supervisor or a colleague ☐
Stopped using that tool for that kind of task ☐ Reported it to someone at
[INSTITUTION NAME] ☐ This has not happened to me ☐ It happened, but I
only noticed later ☐ I would rather not answer

Cells: D4-people · Evidence anchor: The institution's formal failure reporting channel and its incident log. If no respondents check "Reported it to someone," you should investigate whether an accessible reporting channel actually exists. Structured failure-mode reporting (the report §2.4, competency 6) requires a clear institutional mechanism to receive and log those reports.

Section 5: Telling people you used it (4 items)

Q22 ★ ⇌ PAR-DISCLOSE-LAST (core) In the most recent piece of written work you submitted to anyone (advisor or supervisor, milestone panel, journal, funder), did you say that you had used an AI tool?

☐ Yes, in the document itself ☐ Yes, on a form ☐ Yes, verbally only ☐ No ☐ I did not use an AI tool on that piece of work

Cells: D2-process · Evidence anchor: Thesis submission disclosure fields, milestone review paperwork, and disclosure statements on submitted manuscripts. When the British Medical Journal introduced a structured disclosure field, it registered 5.7 percent compliance across 25,114 submissions to 49 BMJ journals in its first seven months (the report §1.1). You should expect a similar pattern here.

Q23 ⇌ PAR-DISCLOSE-CONTENT (core) If you disclosed, what did you record? (tick all that apply)

- ☐ The name of the tool ☐ The version of the tool ☐ The date you used it ☐ What you used it for ☐ Which step of the work it touched ☐ The prompt or prompts you used ☐ A general statement that AI was used, with no detail

Cells: D2-process, D4-people · Evidence anchor: The institutional disclosure template and a sample of completed disclosure forms. The first five options reflect the established standard used by the University of Helsinki's research guidelines: tool name, version, date of use, how the tool was used, and the specific research step involved. the report §1.3 identifies Helsinki as one of six Class D institutions.

Q24 Do you keep a record of your prompts and the tool's outputs?

- ☐ Yes, systematically ☐ Sometimes ☐ No ☐ I did not know this was expected

Cells: D4-people, D3-process · Evidence anchor: Research data management plans submitted for the doctoral program, as well as institutional records governance policies regarding the retention of AI session logs (the report §2.3, auditability).

Q25 Have you ever chosen **not** to say you used an AI tool because you were unsure whether it was allowed, or worried about how it would look?

- ☐ Yes ☐ No ☐ I would rather not answer

Cells: D2-process, D2-people · Evidence anchor: None. This is a composite indication of disclosure reluctance. The item does not distinguish uncertainty about the rule from concern about how disclosure would be received, so investigate those possible explanations separately before choosing a remedy.

Section 6: What you have been taught (6 items)

These six items cover the six foundational research competencies detailed in §2.4 of the report.

For every item in this section:

- ☐ Yes, in a formal course or workshop ☐ Yes, informally (my advisor, supervisor or a colleague) ☐ No ☐ I do not know what this means

Q26 ★ ⇌ PAR-COMP-1 (core) - citation verification Checking every reference an AI tool gives you against the original source before you use it. Has anyone at [INSTITUTION NAME] taught you this? *Cells: D4-people, D4-process · Evidence anchor: Methods curriculum learning outcomes and the library's research skills program. An answer of "I do not know what this means" indicates that the concept was not recognized by the respondent. Report it separately, never merge it into "No", and do not treat it as a direct measure of competence.*

Q27 ★ ⇌ PAR-COMP-2 (core) - model-and-parameter specification Recording which AI model and version you used, on what date, with which settings, so that someone else could repeat what you did. Has anyone at [INSTITUTION NAME] taught you this? *Cells: D4-people, D4-process · Evidence anchor: Methods curriculum outcomes and research data management plan templates.*

Q28 ⇌ PAR-COMP-3 (core) - prompt-as-fork-in-the-garden discipline Deciding your prompts in advance, and testing whether a different wording changes the result. Has anyone at [INSTITUTION NAME] taught you this? *Cells: D4-people, D4-process · Evidence anchor: Methods curriculum outcomes covering preregistration and researcher degrees of freedom, as well as formal preregistration templates that explicitly incorporate prompt design.*

Q29 ⇌ PAR-COMP-4 (core) - model-heterogeneity in adversarial review Using a tool based on a different underlying model family to challenge work you produced with the first one. A different product or company may still use the same model. Has anyone at [INSTITUTION NAME] taught you this? Cells: D4-people, D4-process · Evidence anchor: Methods curriculum learning outcomes. These results should be cross-referenced against Q19, where students report whether they apply this practice.

Q30 ★ ⇌ PAR-COMP-5 (core) - sycophancy detection and human-as-verifier discipline Recognizing that an AI tool agreeing with you is not evidence that you are right, and that you remain the person who checks. Has anyone at [INSTITUTION NAME] taught you this? Cells: D4-people, D4-process · Evidence anchor: Methods curriculum learning outcomes. These results should be read directly alongside responses to Q42.

Q31 ⇌ PAR-COMP-6 (core) - structured failure-mode reporting Reporting how an AI-assisted method can fail, and how often it did, in the same way you report any other method. Has anyone at [INSTITUTION NAME] taught you this? Cells: D4-people, D4-process · Evidence anchor: Methods curriculum learning outcomes. These results should be compared directly with responses to Q21.

Section 7: Your supervisor (4 items)

Q32 ★ ⇌ PAR-SUPERVISION-DISCUSSED (core) Has your advisor or supervisor ever discussed AI use with you?

☐ Never ☐ Once, briefly ☐ More than once ☐ We have discussed it and written down what we agreed

Cells: D1-people, D1-process · Evidence anchor: Records of advisor and supervisor meetings, as well as the official supervision agreement template to see if it includes an explicit AI clause. Compare these responses directly with supervisor answers in A3. If supervisors predominantly report discussing AI "more than once" while students report "never," record a discrepancy between the two anonymous cohort reports. It is not a comparison of perceived with actual supervisory practice.

Q33 ★ ⇌ PAR-WRITTEN-AGREEMENT (core) Do you and your advisor or supervisor have a **written** record of what you agreed about AI use?

☐ Yes ☐ Partial (some of it is written down) ☐ No ☐ I don't know

Cells: D1-process, D1-people · Evidence anchor: Signed supervision agreements held by the graduate school, specifically the count of agreements that contain an AI clause. This item measures baseline demand for instrument G1 (the supervisor-student agreement) and tracks its adoption over time.

Q34 ⇌ PAR-ESCALATION-OWNER If you and your advisor or supervisor disagreed about whether a particular AI use was allowed, do you know who decides?

☐ Yes, and I could name the role or office ☐ I think so, but I am not sure ☐ No ☐ I don't know

Cells: D5-people, D1-people · Evidence anchor: The designated decision-maker specified in institutional policy, such as the research integrity office, the graduate school, or the associate dean for research. While Dimension 5 checks whether these formal roles are defined on paper, this question tests whether that information has actually reached the students who would need to use it.

Q35 ⇌ PAR-SAFE-TO-TELL Would you feel able to tell your advisor or supervisor that you had used an AI tool in a way you were later unsure about?

☐ Yes ☐ Probably ☐ No ☐ I would rather not answer

Cells: D1-people, D2-process · Evidence anchor: None. When researchers fear disclosing their AI use, the institution gets compliant-looking paperwork but zero real insight into practice. These responses should be analyzed alongside Q25.

Section 8: Data you have put into AI tools (4 items)

These questions ask what you report has happened in practice, not what should have happened. As noted in the opening statement, your answers cannot be used against you.

Q36 ★ ⇌ PAR-DATA-ENTERED (core) Which of these have you entered into an AI tool that [INSTITUTION NAME] does **not** provide? (tick all that apply)

☐ My own unpublished data or results ☐ Data about research participants, with names and identifying details removed ☐ Data about research participants that could identify someone ☐ Someone else's unpublished manuscript, thesis or grant application ☐ Material covered by a confidentiality agreement or an industry partner's terms ☐ My own unpublished draft writing ☐ Nothing from this list ☐ I would rather not answer

Cells: D3-systems, D3-process · Evidence anchor: The institutional AI tenant's configuration settings for data residency and training on user inputs (the report §2.3), as well as the documented exception process for research projects that genuinely require consumer-tier tools (the report §4.2, Dimension 3). This item establishes only that the institution did not provide the tool. Compare each reported tool with the approved-tool register and exception records before classifying any use as unapproved or as a data-governance event.

Q37 ★ ⇌ PAR-DATA-CHECKED Before using an AI tool with research data, have you checked where the data is stored and whether the company uses it to train its models?

☐ Always ☐ Sometimes ☐ Never ☐ I did not know I could check ☐ I have not used an AI tool with research data

Cells: D3-people, D3-systems · Evidence anchor: Institutional guidance on approved tools and its distribution logs, along with formal procurement standards governing data residency and the use of inputs for model training (the report §2.3). Report the non-use option separately and exclude it from the applicable denominator.

Q38 ⇌ PAR-ETHICS-COVERS-AI (core) Does your ethics approval or your participants' consent cover having their data processed by an outside AI service?

☐ Yes ☐ No ☐ I don't know ☐ I have not used an outside AI service with personal data ☐ My research does not involve personal data

Cells: D3-process · Evidence anchor: The institutional human research ethics application form, specifically whether it includes an explicit question regarding third-party AI data processing. If the ethics application lacks this question, an answer of "I don't know" reflects an institutional oversight rather than student neglect.

Q39 ⇌ PAR-APPROVED-TOOLS Has anyone at [INSTITUTION NAME] told you which AI tools are approved for unpublished or participant data?

☐ Yes ☐ Partial ☐ No ☐ I don't know

Cells: D3-people, D3-process · Evidence anchor: The official approved tools list, its publication date, and the documented record of how it was distributed to graduate students.

Section 9: Confidence and knowledge (4 items)

The first question asks what you feel confident explaining, while the next two questions test specific knowledge with objective right answers. The main goal here is to compare self-reported confidence against demonstrated understanding.

Q40 ★ ⇌ PAR-CONFIDENCE (core) "I could explain to a new doctoral student how to use AI responsibly in research."

☐ Yes ☐ Partly ☐ No

Cells: D4-people · Evidence anchor: None, because this is intentionally a self-assessment item. Its analytical value comes from comparing these ratings with the self-reported checking practice in Q16 and the knowledge probes in Q18, Q41, and Q42, while keeping those forms of evidence distinct.

Q41 ⇌ PAR-REPRO-PROBE An AI tool helped with part of your analysis. Which of these must you record so that another researcher could repeat what you did? (tick all that apply)

☐ The name and version of the model ☐ The date you used it ☐ The prompt you used ☐ Settings such as temperature or seed, where the tool shows them ☐ Other wordings of the prompt that you tried ☐ None of these are needed if the result is correct

Cells: D4-people, D4-process · Evidence anchor: Methods curriculum learning outcomes related to computational reproducibility. Correct scoring criteria are detailed in the administration notes below.

Q42 ★ ⇌ PAR-CONFIDENT-TONE-PROBE An AI tool gives you an answer in a confident tone, with no hedging. What does the confident tone tell you about whether the answer is correct?

☐ Nothing: a confident tone is not evidence ☐ It makes the answer more likely to be correct ☐ It makes the answer less likely to be correct ☐ It depends which tool it is

Cells: D4-people · Evidence anchor: Methods curriculum outcomes covering sycophancy detection and the requirement for human verification. Correct scoring criteria are detailed in the administration notes below.

Q43 In the past three months, how many times did you decide **not** to use an AI tool for a research task because you judged it was not the right thing to do?

☐ Never ☐ Once or twice ☐ Several times ☐ Often ☐ I have not used AI tools for research

Cells: D1-people · Evidence anchor: None. This is the closest single measure of whether the report's labor-versus-judgment distinction (§2.1) is functioning in daily student practice rather than existing only as a written policy. An answer of "Never" from a student who frequently uses AI is a significant finding.

Section 10: Open comments (2 items, optional)

Q44 (optional, about 80 words) Describe one time an AI tool was confidently wrong in your research work, and what you did next.

Response:
.....
.....

Cells: D4-people · Evidence anchor: Responses should be coded against the institution's taxonomy of research AI failures. These qualitative accounts provide useful input for curriculum design and case material for executive reports.

Q45 (optional, about 40 words) What one thing would help you most with using AI well in your research?

Response:
.....

Cells: D4-process · Evidence anchor: None. Responses should be coded into thematic categories, with the most common priorities informing action planning in instrument A4 (the gap-analysis worksheet).

Administration notes (not for respondents)

Answer key for the three knowledge items

Q18. Correct: *You opened the source and read the passage that supports the claim.* A DOI, a link, or an entry in the tool's source list proves only that a reference was generated, not that the paper actually exists or supports the claim being made. Retrieval augmentation (RAG) often shifts AI errors from completely fabricated citations to real papers that are misattributed or misinterpreted, which is much harder to spot (the report §2.3). Having a second AI tool agree with the first is not independent verification.

Q41. Required options: *The name and version of the model, The date you used it, The prompt you used, and Settings such as temperature or seed, where the tool shows them.* Credited but not strictly required: *Other wordings of the prompt that you tried* (reflecting the discipline of testing whether small prompt changes alter results). Incorrect option: *None of these are needed if the result is correct.* Score the response as correct only if all four required options are selected and the incorrect option is not selected.

Q42. Correct: *Nothing: a confident tone is not evidence.* The other options are incorrect. Confidence of tone is not a valid basis for changing the probability that an answer is correct.

Do not circulate this key with the survey, and do not publish it while the survey is open.

Fielding

- Run A2 and A3 in the exact same time window and close them on the same day. Comparing student and supervisor responses is valid only if both groups are reporting on the same academic period.
- Attempt a census by inviting the full doctoral population. Do not describe the achieved respondent group as representative without comparing its composition with the eligible population.
- Avoid fielding the survey in the two weeks leading up to major thesis submission dates or milestone review deadlines. Keep the survey open for three weeks and send two follow-up reminders.

- Always report the response rate alongside every finding and compare respondents with the eligible population by program stage and broad field where disclosure risk permits. The direction of nonresponse bias cannot be inferred from response rate alone.
- Report every "I don't know" response as its own distinct category. Never combine it with "No", because not knowing the policy represents a communication problem rather than an absence of rules.
- Re-run the survey annually during the same calendar week using identical item wording. Changing a question breaks your longitudinal trend line, so if you need to adjust a topic, retire the old item and add a new one rather than editing an existing question in place.

Classify items before analysis

The cell tags show where an item is conceptually relevant. They are not instructions to average every tagged response into a cell score. Before opening the survey, classify each item by its intended analytic use and record whether a higher response indicates stronger practice, greater exposure, greater risk, or no readiness direction at all.

- **Criterion-aligned self-reported operational indicators:** Q10–Q12, Q14, Q16, Q32–Q33, and Q39. These may test a specific A1 operational claim when the item, population, referent, and time period align. Q20, Q22–Q24, and Q36–Q38 remain contextual unless a local analysis plan defines a directly matching criterion, denominator, referent, and time period before fielding.
- **Training-exposure indicators:** Q26–Q31. Report these as evidence of whether training reached respondents. They do not by themselves demonstrate competence.
- **Knowledge probes:** Q18, Q41, and Q42. Score and report each probe separately. Do not combine them with exposure, self-reported practice, or confidence items as if they formed one scale.
- **Context, use, belief, and qualitative indicators:** all remaining items. Use these to describe the population, interpret direct operational findings, and plan improvements. Do not use them to raise or lower a maturity level by arithmetic.

Only a predeclared, directly aligned operational indicator may enter the criterion-specific checks in A4. Report the response distribution and eligible number for every item used. Never calculate a cell-level mean across unlike items simply because they share a cell tag.

The three gaps this survey exists to examine

Gap 1: The document-to-student implementation gap (A1 against A2). For each directly observable criterion that A1 rates as *established* or *leading*, check whether the aligned student item supports the claim. Keep exposure, knowledge, and contextual items beside that comparison rather than averaging them into it. Important comparisons include written task boundaries against Q12, disclosure operation against Q22 and Q23, and approved-tool communication against Q39.

Gap 2: The student-supervisor cohort gap (A3 against A2). Compare aligned parallel-form items across the two respondent groups. These are anonymous cohort comparisons, not matched student-supervisor pairs. A difference is most directly interpretable when both groups report on the same recipient-facing practice, such as whether AI was discussed or recorded. Questions Q32 and Q33 are the primary indicators here. For policy, system, or supervisor-only facts, report the two perspectives side by side and investigate the reason for the difference.

Gap 3: The confidence pattern (inside A2). Cross-tabulate Q40 with Q16, Q18, Q41, and Q42, while keeping the self-reported practice and knowledge results separate. A respondent who reports confidence but misses a knowledge probe may need different support from someone who knows the rule but does not apply it. Report those patterns without treating the five items as a single confidence or competence score.

How gaps are incorporated into the grid. A2 does not automatically overwrite an A1 documentary level. A4 tests individual operational criteria against predeclared, directly aligned survey evidence

and records any discrepancy. [COMMITTEE] then records the resulting level and its reasoning. Context, training exposure, knowledge, and qualitative responses remain important diagnostic evidence, but they are not pooled into a maturity score. When presenting findings, report the maturity levels across all five dimensions along with the binding axis for each, rather than reducing the institution to a single uninformative score.

Parallel-form register

This register records the topic-link keys available in A2. A3 currently uses a subset of them across twelve supervisor items, and several A3 items refer to more than one A2 key. The **core** marker identifies priority topics for local adaptation. A shared key is not, by itself, evidence that two response distributions can be compared directly.

Key	A2 item	What A3 asks the supervisor	Core
PAR-MODE-SEARCH	Q4	Own use, and belief about students' use	
PAR-MODE-COAUTHOR	Q5	Own use, and belief about students' use	
PAR-MODE-VALIDATOR	Q6	Own use, and belief about students' use	
PAR-MODE-TUTOR	Q7	Own use, and belief about students' use	
PAR-POLICY-KNOWN	Q10	Whether a written policy exists	●
PAR-POLICY-LOCATE	Q11	Name or locate it	●
PAR-TASK-RULES	Q12	Whether they have been told, in writing, which tasks are permitted	●
PAR-THESIS-DECLARE	Q14	Whether AI use must be declared at submission, and what must be declared	●
PAR-CITE-CHECK	Q16	Proportion of AI-supplied references checked at source in their own recent work	●
PAR-FAKE-REF-SEEN	Q17	Whether they have seen a fabricated or misattributed reference	
PAR-CITE-PROBE	Q18	Identical knowledge item, identical options	●
PAR-MODEL-HETERO	Q19	Whether they use a different model family to challenge their own work	●
PAR-DISCLOSE-LAST	Q22	Disclosure on their own most recent submitted work	●
PAR-DISCLOSE-CONTENT	Q23	What they recorded	●
PAR-COMP-1 ... PAR-COMP-6	Q26–Q31	Whether they have been <i>trained</i> on each competency, and whether they teach it	●
PAR-SUPERVISION-DISCUSSED	Q32	Whether they have discussed AI with their doctoral students	●
PAR-WRITTEN-AGREEMENT	Q33	Whether anything is written down	●
PAR-ESCALATION-OWNER	Q34	Who decides a disagreement	
PAR-SAFE-TO-TELL	Q35	Whether they believe a doctoral student would tell them	
PAR-DATA-ENTERED	Q36	What they have entered into non-institutional tools	●

Key	A2 item	What A3 asks the supervisor	Core
PAR-DATA-CHECKED	Q37	Whether they check storage and training-on-input	
PAR-ETHICS-COVERS-AI	Q38	Whether ethics approval covers third-party AI processing	●
PAR-APPROVED-TOOLS	Q39	Whether they know the approved-tool list	
PAR-CONFIDENCE	Q40	"I could explain to a new doctoral student how to use AI responsibly in research."	●
PAR-REPRO-PROBE	Q41	Identical knowledge item	
PAR-CONFIDENT-TONE-PROBE	Q42	Identical knowledge item	

Table 2. Alignment of parallel survey items and core priority designations across instruments.

The supervisor survey (A3) is authored independently. Use the key to locate related items, then record differences in wording, referent, time period, denominator, and response scale before analysis. Report items side by side when those features do not align.

The 24-item short form

The short form includes the 24 items marked with ★: Q1, Q2, Q4, Q5, Q6, Q7, Q8, Q10, Q12, Q13, Q14, Q16, Q18, Q20, Q22, Q26, Q27, Q30, Q32, Q33, Q36, Q37, Q40, and Q42. It takes approximately five minutes to complete. You should use this shorter version when institutional survey fatigue makes the full survey impractical, or during alternate years in an ongoing evaluation cycle. Be sure to document in your final report which version was fielded.

Crosswalk: item to report section

This crosswalk records conceptual relevance. A grid-cell entry does not make an item a scored indicator for that cell. Use the analytic classification above and the criterion-specific rules in A4 before drawing any maturity inference.

Item	Topic	Report section	Grid cells
Q1	Stage of doctoral program	§2.5 (process axis: progression milestones)	D1-process
Q2	Field	Appendix G (discipline-specific extension)	D4-process
Q3	First language	Appendix C, Classes 10 and 11	D3-systems, D4-people
Q4	Search mode use	§2.2 (search)	D2-people, D3-systems
Q5	Co-author mode use	§2.2 (co-author)	D2-people, D2-process
Q6	Validator mode use	§2.2 (validator)	D2-people, D4-people
Q7	Tutor mode use	§2.2 (tutor)	D4-people, D4-process
Q8	Account tier	§2.3 (data residency and training-on-input)	D3-systems, D3-process

Item	Topic	Report section	Grid cells
Q9	Specialized and embedded uses	Appendix C, Classes 5, 7, 9, 10	D3-systems, D2-people
Q10	Policy exists	§1.3, §2.5	D2-people, D1-people
Q11	Policy locatable	§1.3, §2.5	D2-people
Q12	Task-level rules in writing	§2.1, Appendix G	D1-people, D1-process
Q13	Where the rules were learned	§2.4, §4.2	D2-people, D4-process
Q14	Thesis declaration	§2.1, Appendix A (D1 <i>established</i>)	D1-process, D2-process
Q15	Funder and publisher rules	§3.1, §3.3	D2-people, D2-process
Q16	References checked at source	§2.4 (competency 1), §2.2 (search)	D4-people, D2-process
Q17	Fabricated or misattributed reference seen	§2.2 (search), §2.4 (competency 1)	D4-people
Q18	Knowledge probe: verification	§2.4 (competency 1), §2.3	D4-people, D4-process
Q19	Different model family for critique	§2.4 (competency 4)	D4-people
Q20	Used output that could not be checked	§2.2 (co-author)	D1-people, D2-process
Q21	Response to a confident error	§2.4 (competency 6), Appendix C Class 3	D4-people
Q22	Disclosure on last output	§2.2, §3.3, §1.1	D2-process
Q23	Content of the disclosure	§2.2, §2.4 (competency 2), §1.3	D2-process, D4-people
Q24	Prompt and output records	§2.3 (auditability), §2.4 (competency 2)	D4-people, D3-process
Q25	Chose not to disclose	§1.1, §2.2	D2-process, D2-people
Q26	Taught: citation verification	§2.4 (competency 1)	D4-people, D4-process
Q27	Taught: model-and-parameter specification	§2.4 (competency 2)	D4-people, D4-process
Q28	Taught: prompt-as-fork-in-the-garden discipline	§2.4 (competency 3)	D4-people, D4-process
Q29	Taught: model-heterogeneity in adversarial review	§2.4 (competency 4)	D4-people, D4-process
Q30	Taught: sycophancy detection and human-as-verifier discipline	§2.4 (competency 5)	D4-people, D4-process
Q31	Taught: structured failure-mode reporting	§2.4 (competency 6)	D4-people, D4-process

Item	Topic	Report section	Grid cells
Q32	Supervisor discussed AI	§2.1, §4.2 (D1)	D1-people, D1-process
Q33	Written record with supervisor	§2.1, Appendix G (supervisor-approval rows)	D1-process, D1-people
Q34	Who decides a disagreement	§2.5 (people axis)	D1-people
Q35	Safe to tell the supervisor	§2.5 (process axis), §1.2 (transparent)	D1-people, D2-process
Q36	Data entered into non-institutional tools	§2.3 (data residency), §4.2 (Dimension 3)	D3-systems, D3-process
Q37	Checked storage and training-on-input	§2.3	D3-people, D3-systems
Q38	Ethics approval covers third-party AI	§2.3, §2.5 (regulatory pressure points)	D3-process
Q39	Approved-tool list known	§2.3 (procurement gate)	D3-people, D3-process
Q40	Confidence self-report	§2.4	D4-people
Q41	Knowledge probe: reproducibility	§2.4 (competency 2)	D4-people, D4-process
Q42	Knowledge probe: confident tone	§2.4 (competency 5), Appendix C Class 3	D4-people
Q43	Declined to use AI on judgment grounds	§2.1	D1-people
Q44	Open: a confident error	§2.4 (competency 6)	D4-people
Q45	Open: what would help	§5.3, §5.5	D4-process

Table 3. Mapping of individual survey items and topics to report sections and grid cells.

Plain-text version for survey-platform import

Each question block represents a single survey item, and `SHOW_IF` lines specify display and skip logic. Make sure to replace `[INSTITUTION NAME]` and any other bracketed fields before importing. Evidence anchors, cell mappings, ★ marks, and ⇔ parallel keys are omitted here intentionally because respondents must not see them. Keep the `PARKEY` line intact, because it serves as a metadata comment for your analysis file and most survey platforms will import it directly into a question-notes field.

`SURVEY TITLE: AI use in doctoral research at [INSTITUTION NAME]`

`INTRO PAGE: [paste the "Before you begin" and "A note on words" sections here, unedited]`

`SCALE_A = Yes | Partial | No | I don't know`

`SCALE_B = Never | A few times in the past three months | About once a month | About once a week | Most days`

`SCALE_C = All | Most | About half | A few | None | I have not done this`

`PAGE 1 – About you`

[Q1] TYPE=single

TEXT: Where are you in your doctorate?

OPTIONS: First year, before my first formal progress review or milestone | Past that first review, in the middle years | Final year | Submitted, waiting for examination | Other

[Q2] TYPE=single

TEXT: Which best describes your field?

OPTIONS: Humanities | Social sciences | Business, law, economics | Life and health sciences | Physical sciences, engineering, computing | Creative and performing arts | Other

[Q3] TYPE=single

TEXT: Is English your first language?

OPTIONS: Yes | No | One of several first languages

PAGE 2 – Which AI tools you use

PAGE TEXT: "AI tool" means any tool that generates text, code, images, summaries, transcripts, translations or citations for you, whether or not it is sold as artificial intelligence.

[Q4] TYPE=single | SCALE=B | PARKEY=PAR-MODE-SEARCH

TEXT: In the past three months, how often have you used an AI tool to find or scope literature?

[Q5] TYPE=single | SCALE=B | PARKEY=PAR-MODE-COAUTHOR

TEXT: In the past three months, how often have you used an AI tool to draft, rewrite, shorten or polish your own writing?

[Q6] TYPE=single | SCALE=B | PARKEY=PAR-MODE-VALIDATOR

TEXT: In the past three months, how often have you used an AI tool to check, critique or find errors in your own argument, method or results?

[Q7] TYPE=single | SCALE=B | PARKEY=PAR-MODE-TUTOR

TEXT: In the past three months, how often have you used an AI tool to explain a concept, a method or a piece of statistics to you?

[Q8] TYPE=single

TEXT: Which account do you mostly use for research work?

OPTIONS: An account [INSTITUTION NAME] provides | A free personal account | A paid personal account I pay for myself | An account provided by a collaborator or an external partner | A tool that runs on my own computer | I do not use AI tools for research

[Q9] TYPE=multi

TEXT: Which of these have you used for research work in the past three months? Tick all that apply.

OPTIONS: Writing or fixing code, or building an analysis pipeline | Screening titles and abstracts for a review | Transcribing interviews, focus groups or fieldnotes | Translating text | Making images or diagrams for talks or teaching | Making images or figures that present data | AI features inside tools I do not think of as AI tools (search-engine answers, reference-manager summaries, meeting transcription, writing assistants in email or word processors) | None of these

EXCLUSIVE: None of these

PAGE 3 – What you have been told

[Q10] TYPE=single | SCALE=A | PARKEY=PAR-POLICY-KNOWN

TEXT: Does [INSTITUTION NAME] have a written rule or policy about using AI in doctoral research?

[Q11] TYPE=text_short | REQUIRED=no | SHOW_IF=Q10 IN (Yes, Partial) | PARKEY=PAR-POLICY-LOCATE

TEXT: What is it called, or where is it? A title, a web address, or a description such as "a page on the graduate school site" is enough. Write "cannot find it" if that is the honest answer.

[Q12] TYPE=single | SCALE=A | PARKEY=PAR-TASK-RULES

TEXT: Have you been told in writing which research tasks you may use AI for and which you may not?

[Q13] TYPE=multi

TEXT: Where did you learn what the rules are? Tick all that apply.

OPTIONS: Induction or orientation | A methods or research-training course | My advisor or supervisor | The thesis submission form | An email or a web page from my school or college, or from the graduate school | Other doctoral students | The AI tool's own terms or warnings | A journal or funder, not my institution | I have not learned them

EXCLUSIVE: I have not learned them

[Q14] TYPE=single | PARKEY=PAR-THESIS-DECLARE

TEXT: Do you know whether you must declare AI use when you submit your thesis?

OPTIONS: Yes, and I know what I must declare | Yes, but I do not know what I must declare | No, there is no requirement | I don't know

[Q15] TYPE=single

TEXT: Do you know what the rules are for using AI in grant or scholarship applications, and in manuscripts you submit to journals?

OPTIONS: Yes, both | One but not the other | No | I don't know | Not relevant to me yet

PAGE 4 – Checking what the tool gives you

[Q16] TYPE=single | SCALE=C | PARKEY=PAR-CITE-CHECK

TEXT: In your most recent literature search that used an AI tool: how many of the references it gave you did you open and read at the original source?

[Q17] TYPE=single | PARKEY=PAR-FAKE-REF-SEEN

TEXT: Have you ever found that a reference an AI tool gave you did not exist, or did not say what the tool claimed it said?

OPTIONS: Yes | No | I have not checked closely enough to know

[Q18] TYPE=single | PARKEY=PAR-CITE-PROBE

TEXT: Which one of these confirms that a reference is real and says what an AI tool claims it says?

OPTIONS: The tool gave a DOI or a link | The reference appears in the tool's list of

sources | You opened the source and read the passage that supports the claim | A second AI tool agreed that the reference is correct

RANDOMIZE: yes

[Q19] TYPE=single | PARKEY=PAR-MODEL-HETERO

TEXT: When you use an AI tool to challenge or check your own argument, do you use a tool based on a different underlying model family from the one that helped you produce the argument? A different product or company may still use the same underlying model.

OPTIONS: Yes, always | Sometimes | No | I had not thought about this | I do not use AI to check my arguments

[Q20] TYPE=single

TEXT: Have you ever included AI output in a thesis chapter, a manuscript or an application when you could not check it yourself?

OPTIONS: Yes | Partial - I checked some of it | No | I would rather not answer

[Q21] TYPE=multi

TEXT: The last time an AI tool gave you an answer that sounded confident and turned out to be wrong, what did you do? Tick all that apply.

OPTIONS: Checked it against the source and corrected it | Asked the same tool again | Asked a different tool | Asked my advisor, supervisor or a colleague | Stopped using that tool for that kind of task | Reported it to someone at [INSTITUTION NAME] | This has not happened to me | It happened, but I only noticed later | I would rather not answer

EXCLUSIVE: This has not happened to me, I would rather not answer

PAGE 5 - Telling people you used it

[Q22] TYPE=single | PARKEY=PAR-DISCLOSE-LAST

TEXT: In the most recent piece of written work you submitted to anyone - advisor or supervisor, milestone panel, journal, funder - did you say that you had used an AI tool?

OPTIONS: Yes, in the document itself | Yes, on a form | Yes, verbally only | No | I did not use an AI tool on that piece of work

[Q23] TYPE=multi | SHOW_IF=Q22 IN ("Yes, in the document itself", "Yes, on a form", "Yes, verbally only") | PARKEY=PAR-DISCLOSE-CONTENT

TEXT: What did you record? Tick all that apply.

OPTIONS: The name of the tool | The version of the tool | The date you used it | What you used it for | Which step of the work it touched | The prompt or prompts you used | A general statement that AI was used, with no detail

[Q24] TYPE=single

TEXT: Do you keep a record of your prompts and the tool's outputs?

OPTIONS: Yes, systematically | Sometimes | No | I did not know this was expected

[Q25] TYPE=single

TEXT: Have you ever chosen not to say you used an AI tool because you were unsure whether it was allowed, or worried about how it would look?

OPTIONS: Yes | No | I would rather not answer

PAGE 6 - What you have been taught

PAGE TEXT: For each one, choose the option that fits best.

GRID OPTIONS (all six rows): Yes, in a formal course or workshop | Yes, informally – my advisor, supervisor or a colleague | No | I do not know what this means

[Q26] TYPE=grid_row | PARKEY=PAR-COMP-1

TEXT: Citation verification – checking every reference an AI tool gives you against the original source before you use it. Has anyone at [INSTITUTION NAME] taught you this?

[Q27] TYPE=grid_row | PARKEY=PAR-COMP-2

TEXT: Model-and-parameter specification – recording which AI model and version you used, on what date, with which settings, so that someone else could repeat what you did. Has anyone at [INSTITUTION NAME] taught you this?

[Q28] TYPE=grid_row | PARKEY=PAR-COMP-3

TEXT: Prompt-as-fork-in-the-garden discipline – deciding your prompts in advance, and testing whether a different wording changes the result. Has anyone at [INSTITUTION NAME] taught you this?

[Q29] TYPE=grid_row | PARKEY=PAR-COMP-4

TEXT: Model-heterogeneity in adversarial review – using a tool based on a different underlying model family to challenge work you produced with the first one. A different product or company may still use the same model. Has anyone at [INSTITUTION NAME] taught you this?

[Q30] TYPE=grid_row | PARKEY=PAR-COMP-5

TEXT: Sycophancy detection and human-as-verifier discipline – recognizing that an AI tool agreeing with you is not evidence that you are right, and that you remain the person who checks. Has anyone at [INSTITUTION NAME] taught you this?

[Q31] TYPE=grid_row | PARKEY=PAR-COMP-6

TEXT: Structured failure-mode reporting – reporting how an AI-assisted method can fail, and how often it did, in the same way you report any other method. Has anyone at [INSTITUTION NAME] taught you this?

PAGE 7 – Your supervisor

[Q32] TYPE=single | PARKEY=PAR-SUPERVISION-DISCUSSED

TEXT: Has your advisor or supervisor ever discussed AI use with you?

OPTIONS: Never | Once, briefly | More than once | We have discussed it and written down what we agreed

[Q33] TYPE=single | PARKEY=PAR-WRITTEN-AGREEMENT

TEXT: Do you and your advisor or supervisor have a written record of what you agreed about AI use?

OPTIONS: Yes | Partial – some of it is written down | No | I don't know

[Q34] TYPE=single | PARKEY=PAR-ESCALATION-OWNER

TEXT: If you and your advisor or supervisor disagreed about whether a particular AI use was allowed, do you know who decides?

OPTIONS: Yes, and I could name the role or office | I think so, but I am not sure | No | I don't know

[Q35] TYPE=single | PARKEY=PAR-SAFE-TO-TELL

TEXT: Would you feel able to tell your advisor or supervisor that you had used an AI tool in a way you were later unsure about?

OPTIONS: Yes | Probably | No | I would rather not answer

PAGE 8 – Data you have put into AI tools

PAGE TEXT: These questions ask what you report has happened, not what should have happened. Your answers cannot be used against you – see the statement at the start.

[Q36] TYPE=multi | PARKEY=PAR-DATA-ENTERED

TEXT: Which of these have you entered into an AI tool that [INSTITUTION NAME] does not provide? Tick all that apply.

OPTIONS: My own unpublished data or results | Data about research participants, with names and identifying details removed | Data about research participants that could identify someone | Someone else's unpublished manuscript, thesis or grant application | Material covered by a confidentiality agreement or an industry partner's terms | My own unpublished draft writing | Nothing from this list | I would rather not answer

EXCLUSIVE: Nothing from this list, I would rather not answer

[Q37] TYPE=single | PARKEY=PAR-DATA-CHECKED

TEXT: Before using an AI tool with research data, have you checked where the data is stored and whether the company uses it to train its models?

OPTIONS: Always | Sometimes | Never | I did not know I could check

[Q38] TYPE=single | PARKEY=PAR-ETHICS-COVERS-AI

TEXT: Does your ethics approval or your participants' consent cover having their data processed by an outside AI service?

OPTIONS: Yes | No | I don't know | I have not used an outside AI service with personal data | My research does not involve personal data

[Q39] TYPE=single | SCALE=A | PARKEY=PAR-APPROVED-TOOLS

TEXT: Has anyone at [INSTITUTION NAME] told you which AI tools are approved for unpublished or participant data?

PAGE 9 – Confidence and knowledge

[Q40] TYPE=single | PARKEY=PAR-CONFIDENCE

TEXT: "I could explain to a new doctoral student how to use AI responsibly in research."

OPTIONS: Yes | Partly | No

[Q41] TYPE=multi | PARKEY=PAR-REPRO-PROBE

TEXT: An AI tool helped with part of your analysis. Which of these must you record so that another researcher could repeat what you did? Tick all that apply.

OPTIONS: The name and version of the model | The date you used it | The prompt you used | Settings such as temperature or seed, where the tool shows them | Other wordings of the prompt that you tried | None of these are needed if the result is correct

EXCLUSIVE: None of these are needed if the result is correct

[Q42] TYPE=single | PARKEY=PAR-CONFIDENT-TONE-PROBE

TEXT: An AI tool gives you an answer in a confident tone, with no hedging. What does the confident tone tell you about whether the answer is correct?

OPTIONS: Nothing – a confident tone is not evidence | It makes the answer more likely to be correct | It makes the answer less likely to be correct | It depends which tool it is
RANDOMIZE: yes

[Q43] TYPE=single

TEXT: In the past three months, how many times did you decide not to use an AI tool for a research task because you judged it was not the right thing to do?

OPTIONS: Never | Once or twice | Several times | Often | I have not used AI tools for research

PAGE 10 – Open comments (optional)

[Q44] TYPE=text_long | REQUIRED=no | MAXWORDS=80

TEXT: Describe one time an AI tool was confidently wrong in your research work, and what you did next.

[Q45] TYPE=text_long | REQUIRED=no | MAXWORDS=40

TEXT: What one thing would help you most with using AI well in your research?

CLOSING PAGE: Thank you. A summary of the results goes to [COMMITTEE] on [DATE] and the same summary goes to every doctoral student. If you want to raise something this survey did not cover, contact [ROLE] at [EMAIL].

instats

Michael J. Zyphur, PhD · Professor and Director, Instats · instats.org · support@instats.org

Part of the Instats AI Readiness Pack · instats.org/publications/ai-readiness · CC BY 4.0 - adapt freely for institutional use, with attribution. Companion to Responsible AI in Academic Research: A Competency Framework for Research Training (DOI 10.61700/t31oy23grr).